

2_Diabetes_Predictions_Using_SVM

June 25, 2024

0.1 Diabetes Predictions Using SVM - *Vignesh Prabhu*

Diabetes prediction using the PIMA dataset involves leveraging a **Support Vector Machine** (SVM) model in machine learning. SVM, a powerful classification algorithm, analyzes key health indicators from the dataset (like glucose levels, BMI, etc.) to predict the likelihood of diabetes in individuals. This approach aims to provide accurate predictions based on historical data, aiding early detection and proactive management of the disease.

Importing the Dependencies

```
[37]: import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import accuracy_score
```

Data Collection And Pre-Processing

```
[38]: # Loading The dataset to pandas dataframe
diabetes=pd.read_csv('/content/diabetes.csv')
```

```
[39]: #To display the first 5 Data from dataframe

diabetes.head()
```

```
[39]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|---|-------------|---------|---------------|---------------|---------|------|---|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |
| 3 | 0.167 | 21 | 0 |
| 4 | 2.288 | 33 | 1 |

```
[40]: #To check no.of Rows and Columns
diabetes.shape
```

```
[40]: (768, 9)
```

```
[41]: #Statistical Measures of Data
diabetes.describe()
```

```
[41]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin \ |
|-------|-------------|------------|---------------|---------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 |

| | BMI | DiabetesPedigreeFunction | Age | Outcome |
|-------|------------|--------------------------|------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 31.992578 | 0.471876 | 33.240885 | 0.348958 |
| std | 7.884160 | 0.331329 | 11.760232 | 0.476951 |
| min | 0.000000 | 0.078000 | 21.000000 | 0.000000 |
| 25% | 27.300000 | 0.243750 | 24.000000 | 0.000000 |
| 50% | 32.000000 | 0.372500 | 29.000000 | 0.000000 |
| 75% | 36.600000 | 0.626250 | 41.000000 | 1.000000 |
| max | 67.100000 | 2.420000 | 81.000000 | 1.000000 |

```
[42]: diabetes['Outcome'].value_counts() #0 means non-diabetic and 1 means diabetic
```

```
[42]: Outcome
0    500
1    268
Name: count, dtype: int64
```

```
[43]: diabetes.groupby('Outcome').mean()
```

```
[43]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin \ |
|---------|-------------|------------|---------------|---------------|------------|
| Outcome | | | | | |
| 0 | 3.298000 | 109.980000 | 68.184000 | 19.664000 | 68.792000 |
| 1 | 4.865672 | 141.257463 | 70.824627 | 22.164179 | 100.335821 |

| | BMI | DiabetesPedigreeFunction | Age |
|---------|-----------|--------------------------|-----------|
| Outcome | | | |
| 0 | 30.304200 | 0.429734 | 31.190000 |
| 1 | 35.142537 | 0.550500 | 37.067164 |

Separate Data and Label

```
[44]: X=diabetes.drop(columns='Outcome', axis=1) #axis =1 for column
      Y=diabetes['Outcome'] #Label
```

```
[45]: print(X)
      print(Y)
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | \ |
|-----|-------------|---------|---------------|---------------|---------|------|---|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | |
| .. | ... | ... | ... | ... | ... | ... | |
| 763 | 10 | 101 | 76 | 48 | 180 | 32.9 | |
| 764 | 2 | 122 | 70 | 27 | 0 | 36.8 | |
| 765 | 5 | 121 | 72 | 23 | 112 | 26.2 | |
| 766 | 1 | 126 | 60 | 0 | 0 | 30.1 | |
| 767 | 1 | 93 | 70 | 31 | 0 | 30.4 | |

| | DiabetesPedigreeFunction | Age |
|-----|--------------------------|-----|
| 0 | 0.627 | 50 |
| 1 | 0.351 | 31 |
| 2 | 0.672 | 32 |
| 3 | 0.167 | 21 |
| 4 | 2.288 | 33 |
| .. | ... | ... |
| 763 | 0.171 | 63 |
| 764 | 0.340 | 27 |
| 765 | 0.245 | 30 |
| 766 | 0.349 | 47 |
| 767 | 0.315 | 23 |

[768 rows x 8 columns]

| | |
|-----|---|
| 0 | 1 |
| 1 | 0 |
| 2 | 1 |
| 3 | 0 |
| 4 | 1 |
| .. | |
| 763 | 0 |
| 764 | 0 |
| 765 | 0 |
| 766 | 1 |
| 767 | 0 |

Name: Outcome, Length: 768, dtype: int64

0.2 Data Standardization

```
[46]: scaler=StandardScaler()
```

```
[47]: scaler.fit(X) # Fitting inconsistent data
```

```
[47]: StandardScaler()
```

```
[48]: Standardized_data=scaler.transform(X) #Transform Those data
```

```
[49]: print(Standardized_data) #Standardized data
```

```
[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
   1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
  -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
  -0.10558415]
 ...
 [ 0.3429808   0.00330087  0.14964075 ... -0.73518964 -0.68519336
  -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
   1.17073215]
 [-0.84488505 -0.8730192   0.04624525 ... -0.20212881 -0.47378505
  -0.87137393]]
```

```
[50]: X=Standardized_data
      Y=diabetes['Outcome']
```

```
[51]: print(X)
      print(Y)
```

```
[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
   1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
  -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
  -0.10558415]
 ...
 [ 0.3429808   0.00330087  0.14964075 ... -0.73518964 -0.68519336
  -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
   1.17073215]
 [-0.84488505 -0.8730192   0.04624525 ... -0.20212881 -0.47378505
  -0.87137393]]
0      1
1      0
2      1
```

```

3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64

```

Train Test Split

```
[52]: X_train, X_test, Y_train, Y_test=train_test_split (X,Y, test_size=0.2,
↳stratify=Y, random_state=2)
```

```
[53]: print(X.shape, X_train.shape, X_test.shape)
```

```
(768, 8) (614, 8) (154, 8)
```

Model Train

```
[54]: classifier=svm.SVC(kernel='linear') #Support Vector Machine Classifier
```

```
[55]: #training the SVM classifier
classifier.fit(X_train, Y_train) #training the model with training data
```

```
[55]: SVC(kernel='linear')
```

Model Evaluation

```
[56]: #Accuracy Score on the training data
X_train_prediction=classifier.predict(X_train)
training_data_accuracy=accuracy_score(X_train_prediction, Y_train)
```

```
[57]: print('Accuracy score of the training data : ', training_data_accuracy) #
↳accuracy more than 75 is fine
```

```
Accuracy score of the training data : 0.7866449511400652
```

```
[58]: #Accuracy score on the test data
Y_test_prediction=classifier.predict(X_test)
test_data_accuracy=accuracy_score(Y_test_prediction, Y_test)
```

```
[59]: print('Accuracy score of the test data : ', test_data_accuracy) # accuracy more
↳than 75 is fine
```

```
Accuracy score of the test data : 0.7727272727272727
```

Making Predictive System

```
[60]: input_data=(5,166,72,19,175,25.8,0.587,51)

#changing the data to np array
input_data_as_numpy_array=np.asarray(input_data)

#reshape the array as we are predicting for one instance
input_data_reshaped=input_data_as_numpy_array.reshape(1,-1)

#standardize the input data
std_data=scaler.transform(input_data_reshaped)
print(std_data)

prediction=classifier.predict(std_data)
print(prediction)

if (prediction[0]==0):
    print('The person is not diabetic')
else:
    print('The person is diabetic')
```

[[0.3429808 1.41167241 0.14964075 -0.09637905 0.82661621 -0.78595734
 0.34768723 1.51108316]]

[1]
The person is diabetic

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(